

Architectures, Regularization and Optimization (Part II)

Steve Avsec

Illinois Institute of Technology

March 13, 2025

Overview

- 1 More Deep Learning Regularization
- 2 Optimization
- 3 More Architectures
- 4 Transformers
- 5 A Detour

Bagging

- Idea: Similar to bootstrapping, sample k different training datasets by sampling from the original training data (with replacement).

Bagging

- Idea: Similar to bootstrapping, sample k different training datasets by sampling from the original training data (with replacement).
- Train model j on training set j .

Bagging

- Idea: Similar to bootstrapping, sample k different training datasets by sampling from the original training data (with replacement).
- Train model j on training set j .
- Ensemble the k resulting models by "voting" on the best outcome.

Bagging

- Idea: Similar to bootstrapping, sample k different training datasets by sampling from the original training data (with replacement).
- Train model j on training set j .
- Ensemble the k resulting models by "voting" on the best outcome.
- Downside: Computationally expensive for large networks.

Dropout

- Sample a minibatch from the training data and select a *mask* of the input features + hidden units.

Dropout

- Sample a minibatch from the training data and select a *mask* of the input features + hidden units.
- "Typical" selection is an input feature is included with probability 0.8 and hidden unit with probability 0.5.

Dropout

- Sample a minibatch from the training data and select a *mask* of the input features + hidden units.
- "Typical" selection is an input feature is included with probability 0.8 and hidden unit with probability 0.5.
- Parameters are saved and shared among each iteration.

Dropout

- Sample a minibatch from the training data and select a *mask* of the input features + hidden units.
- "Typical" selection is an input feature is included with probability 0.8 and hidden unit with probability 0.5.
- Parameters are saved and shared among each iteration.
- Even a relatively few draws of masks and few trainings can yield very good models.

Stochastic Gradient Descent

- 1 Select a sequence of learning rates (ε_k) and initial parameter set θ .

Stochastic Gradient Descent

- 1 Select a sequence of learning rates (ε_k) and initial parameter set θ .
- 2 Sample a small batch of the training data.

Stochastic Gradient Descent

- 1 Select a sequence of learning rates (ε_k) and initial parameter set θ .
- 2 Sample a small batch of the training data.
- 3 Compute

$$\mathbf{g} = \frac{1}{m} \nabla L(f(\mathbf{x}, \theta), y)$$

Stochastic Gradient Descent

- 1 Select a sequence of learning rates (ε_k) and initial parameter set θ .
- 2 Sample a small batch of the training data.

- 3 Compute

$$\mathbf{g} = \frac{1}{m} \nabla L(f(\mathbf{x}, \theta), y)$$

- 4 k th iteration,

$$\theta_{k+1} = \theta_k - \varepsilon_k \mathbf{g}$$

Stochastic Gradient Descent

- 1 Select a sequence of learning rates (ε_k) and initial parameter set θ .
- 2 Sample a small batch of the training data.
- 3 Compute

$$\mathbf{g} = \frac{1}{m} \nabla L(f(\mathbf{x}, \theta), y)$$

- 4 k th iteration,

$$\theta_{k+1} = \theta_k - \varepsilon_k \mathbf{g}$$

Notes: Converges if $\sum_{j=1}^{\infty} \varepsilon_j < \infty$, but tends to converge slowly (in high dimensions).

SGD with Momentum

Update SGD with the following changes:

- Choose an initial “velocity” vector $\mathbf{v}$.

SGD with Momentum

Update SGD with the following changes:

- Choose an initial “velocity” vector $\mathbf{v}$.
- Choose learning rate ε and momentum parameter α .

SGD with Momentum

Update SGD with the following changes:

- Choose an initial “velocity” vector $\mathbf{v}$.
- Choose learning rate ε and momentum parameter α .
- Update $\mathbf{v}$ by

$$\mathbf{v}_{k+1} = \alpha \mathbf{v}_k - \varepsilon \mathbf{g}$$

SGD with Momentum

Update SGD with the following changes:

- Choose an initial "velocity" vector $\mathbf{v}$.
- Choose learning rate ε and momentum parameter α .
- Update $\mathbf{v}$ by

$$\mathbf{v}_{k+1} = \alpha \mathbf{v}_k - \varepsilon \mathbf{g}$$

- Update θ by

$$\theta_{k+1} = \theta_k + \mathbf{v}_{k+1}$$

Adaptive Learning Rates

- AdaGrad
- RMSProp
- Adam

Goal: Keep the gradients out of local minima and exploring the parameter space efficiently.

Conjugate Gradient Descent

- Similar to gradient descent:

$$x_{k+1} = x_k + \alpha \nabla f(x_k)$$

Conjugate Gradient Descent

- Similar to gradient descent:

$$x_{k+1} = x_k + \alpha \nabla f(x_k)$$

- Except replace

$$\alpha \nabla f(x_k)$$

with

$$\sum_{j=1}^n \frac{\langle \nabla f(x_k), d_j \rangle}{\langle d_j, Hf(x_k) d_j \rangle}$$

Autoencoders

These are “simple” two-layer neural networks meant to approximate the identity function.

Autoencoders

These are “simple” two-layer neural networks meant to approximate the identity function.

$$e : \mathbb{R}^N \rightarrow \mathbb{R}^n$$

and

$$d : \mathbb{R}^n \rightarrow \mathbb{R}^N$$

such that $id = d \circ e$

Regularization

If we just use

$$L(\mathbf{x}, d, e) = \|\mathbf{x} - d \circ e(\mathbf{x})\|_2,$$

we are quite constrained what we can "learn".

Regularization

If we just use

$$L(\mathbf{x}, d, e) = \|\mathbf{x} - d \circ e(\mathbf{x})\|_2,$$

we are quite constrained what we can "learn".

We can add a regularization constraint

$$L(\mathbf{x}) = \|\mathbf{x} - d \circ e(\mathbf{x})\|_p + \Omega(h)$$

Regularization

If we just use

$$L(\mathbf{x}, d, e) = \|\mathbf{x} - d \circ e(\mathbf{x})\|_2,$$

we are quite constrained what we can “learn”.

We can add a regularization constraint

$$L(\mathbf{x}) = \|\mathbf{x} - d \circ e(\mathbf{x})\|_p + \Omega(h)$$

This can allow us to capture more nuanced information even allowing $n > N$.

Back to RNNs

Last week, we saw that an RNN can be used to map a sequence to a single output (or with a small generalization to a vector of fixed length.)

Back to RNNs

Last week, we saw that an RNN can be used to map a sequence to a single output (or with a small generalization to a vector of fixed length.)

We can construct an RNN that takes a fixed vector and gives a sequence of arbitrary length.

- $\mathbf{a}_t = \mathbf{b} + W\mathbf{h}_{t-1} + U\mathbf{x}$
- $\mathbf{h}_t = g(\mathbf{a}_t)$
- $\mathbf{o}_t = \mathbf{c} + V\mathbf{h}_t$
- $\mathbf{y}_t = d(\mathbf{o}_t)$

Back to RNNs

Last week, we saw that an RNN can be used to map a sequence to a single output (or with a small generalization to a vector of fixed length.)

We can construct an RNN that takes a fixed vector and gives a sequence of arbitrary length.

- $\mathbf{a}_t = \mathbf{b} + W\mathbf{h}_{t-1} + U\mathbf{x}$
- $\mathbf{h}_t = g(\mathbf{a}_t)$
- $\mathbf{o}_t = \mathbf{c} + V\mathbf{h}_t$
- $\mathbf{y}_t = d(\mathbf{o}_t)$

We can combine these to produce *encoder-decoder* sequence-to-sequence architectures.

Context

The loss function in this case usually takes the form

$$L((y_j), (x_k)) = \log(P((y_j), (x_k)))$$

Context

The loss function in this case usually takes the form

$$L((y_j), (x_k)) = \log(P((y_j), (x_k)))$$

The hidden vector **C** which is the output of the encoder/input to the decoder is commonly called the *context*.

Context

The loss function in this case usually takes the form

$$L((y_j), (x_k)) = \log(P((y_j), (x_k)))$$

The hidden vector **C** which is the output of the encoder/input to the decoder is commonly called the *context*.

Having a fixed length vector as context has the drawback that arbitrary length sequences cannot be correctly encoded by fixed length

Attention

To work around this limitation of a fixed context length,
Bahdanau et al. (2015) introduced an attention mechanism.

Attention

To work around this limitation of a fixed context length, Badanau et al. (2015) introduced an attention mechanism. The input sequence (x_k) is mapped to a hidden sequence (h_l) of "annotations".

Attention

To work around this limitation of a fixed context length, Bahdanau et al. (2015) introduced an attention mechanism. The input sequence (x_k) is mapped to a hidden sequence (h_l) of "annotations".

The context vector $\mathbf{c}_j$ is "learned" as a weighted average

$$\mathbf{c}_j = \sum_l \alpha_{j,l} h_l$$

where the weights $\alpha_{j,l}$ are learned from a hidden state s_j where

$$s_j = f(y_{j-1}, s_{j-1}, \mathbf{c}_j)$$

and

$$\alpha_{j,l} = \text{softmax}(a(s_{j-1}, h_l))$$

Attention Continued

The model a is a feedforward model called "alignment".

Attention Continued

The model a is a feedforward model called "alignment".

The hidden units h_l are produced from a bidirectional RNN to get context from both the forward sequence and backward.

Attention Continued

The model a is a feedforward model called "alignment".

The hidden units h_l are produced from a bidirectional RNN to get context from both the forward sequence and backward.

If this sounds complicated and hard to train, Google's AI research team thought so too.

High-Level

Transformers still work on an encoder-decoder type of premise, but use some clever architecture to avoid recurrent units.

High-Level

Transformers still work on an encoder-decoder type of premise, but use some clever architecture to avoid recurrent units.

Encoder is 6 layers with two sub-layers each.

High-Level

Transformers still work on an encoder-decoder type of premise, but use some clever architecture to avoid recurrent units.

Encoder is 6 layers with two sub-layers each.

Decoder is 6 layers of three sub-layers each.

Multi-Head Attention

Define an *attention* function by

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^t}{\sqrt{d_k}} V\right)$$

Multi-Head Attention

Define an *attention* function by

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^t}{\sqrt{d_k}}\right)V$$

And then a multi-head attention layer by

$$\text{MultiHead}(Q, K, V) = \text{Concat}(h_1, h_2, \dots, h_l)W^O$$

where $h_j = \text{Attention}(QW_j^Q, KW_j^K, VW_j^V)$ and W_j^Q, W_j^K, W_j^V are projection matrices.

The Wedding Problem

Suppose there are N people attending a wedding.

The Wedding Problem

Suppose there are N people attending a wedding.

Let $g_{i,j}$ be a random variable denoting that Person i and Person j get along (or not).

The Wedding Problem

Suppose there are N people attending a wedding.

Let $g_{i,j}$ be a random variable denoting that Person i and Person j get along (or not).

Suppose

$$P(g_{i,j} = 1) = P(g_{i,j} = -1) = \frac{1}{2}$$

and that

$$g_{i,j} = g_{j,i}$$

The Wedding Problem

Suppose there are N people attending a wedding.

Let $g_{i,j}$ be a random variable denoting that Person i and Person j get along (or not).

Suppose

$$P(g_{i,j} = 1) = P(g_{i,j} = -1) = \frac{1}{2}$$

and that

$$g_{i,j} = g_{j,i}$$

Suppose everyone has to be seated at one of two tables for dinner.

A Hamiltonian

Let $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_N)$ denote a seating arrangement each $\sigma_j \in \{-1, 1\}$.

A Hamiltonian

Let $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_N)$ denote a seating arrangement each $\sigma_j \in \{-1, 1\}$.

We can look to minimize

$$H(\sigma) = -\frac{1}{\sqrt{N}} \sum_{i < j} \sigma_j \sigma_i g_{i,j}$$

A Hamiltonian

Let $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_N)$ denote a seating arrangement each $\sigma_j \in \{-1, 1\}$.

We can look to minimize

$$H(\sigma) = -\frac{1}{\sqrt{N}} \sum_{i < j} \sigma_j \sigma_i g_{i,j}$$

Gibbs measure: A probability measure on all configurations given by

$$G(\sigma) = \frac{e^{-\beta H(\sigma)}}{Z}$$

where

$$Z = \sum e^{-\beta H(\sigma)}$$

is taken over all configurations.

Variational Autoencoders

Idea: Learn a non-linear function g such that if $\mathbf{z}$ is pulled from some latent known probability distribution, $p(\mathbf{x}; g(\mathbf{z}))$ accurately describes the distribution of $\mathbf{x}$.

Variational Autoencoders

Idea: Learn a non-linear function g such that if $\mathbf{z}$ is pulled from some latent known probability distribution, $p(\mathbf{x}; g(\mathbf{z}))$ accurately describes the distribution of $\mathbf{x}$.
Encoder is given by learning a distribution

$$q(\mathbf{z}|\mathbf{x})$$

and the decoder is just the distribution described above:

$$p(\mathbf{x}|\mathbf{z}) = p(\mathbf{x}; g(\mathbf{z}))$$

Variational Autoencoders

Idea: Learn a non-linear function g such that if $\mathbf{z}$ is pulled from some latent known probability distribution, $p(\mathbf{x}; g(\mathbf{z}))$ accurately describes the distribution of $\mathbf{x}$.
Encoder is given by learning a distribution

$$q(\mathbf{z}|\mathbf{x})$$

and the decoder is just the distribution described above:

$$p(\mathbf{x}|\mathbf{z}) = p(\mathbf{x}; g(\mathbf{z}))$$

Loss function:

$$\mathcal{L}(q) = E[\log(p(x|z))] - D_{KL}(q)||p(z)$$